

# Kubernetes 101

Vinayak Shinde  
CKA | MTS @Portworx



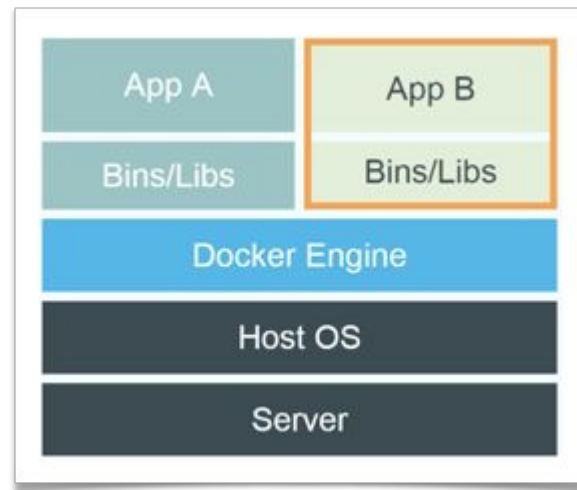
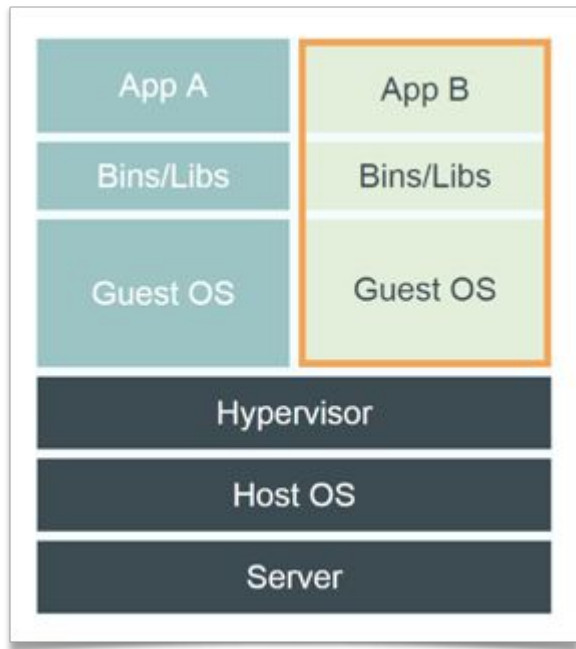
kubernetes

# Agenda

1. What are Containers ?
2. What is Kubernetes ?
3. Key Concepts
4. Architecture
5. Summery



# What are containers ?



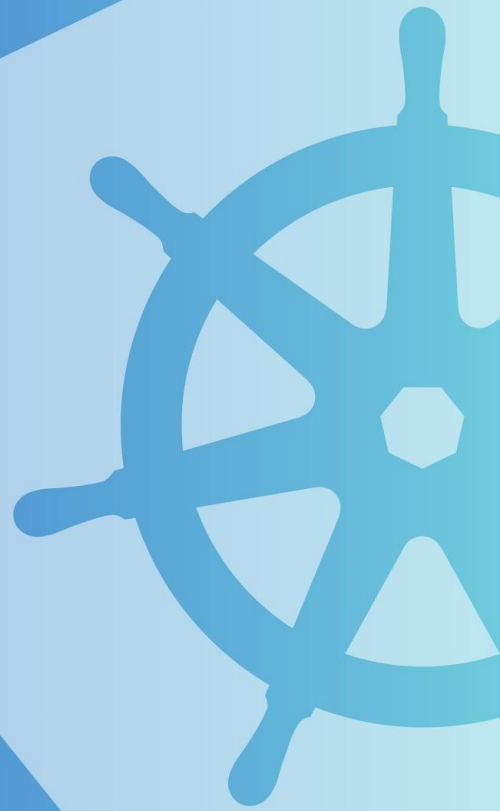
kubernetes

# What is Kubernetes ?

- **A Production-Grade Container Orchestration System** Google-grown, based on Borg and Omega, systems that run inside of Google right now and are proven to work at Google for over 10 years.
- Google spawns billions of containers per week with these systems.
- Created by three Google employees initially during the summer of 2014; grew exponentially and became the first project to get donated to the CNCF.




# Key Concepts




# 1. Pod

- **Atomic unit** or smallest “*unit of work*” of Kubernetes.
- **One or MORE containers** that share
  - Storage: Have access to shared volumes.
  - Network: An IP address and port space, and can find each other via **localhost**.
- **Ephemeral**



# 1. Pod



# Pod Example




```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  labels:
    app: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    resources: {}
```





## 2. Labels

- key-value pairs that are used to identify, describe and group together related sets of objects or resources.
- **NOT** characteristic of uniqueness.
- Have a strict syntax with a slightly limited character set\*.



# 3. Selectors



- Selectors use labels to filter or select objects, and are used throughout Kubernetes.

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-label-example
  labels:
    app: nginx
    env: prod
spec:
  containers:
  - name: nginx
    image: nginx:stable-alpine
    ports:
    - containerPort: 80
  nodeSelector:
    gpu: nvidia
```



### 3. Selector Types



**Equality based** selectors allow for simple filtering (`=`, `==`, or `!=`).

```
selector:  
  matchLabels:  
    gpu: nvidia
```

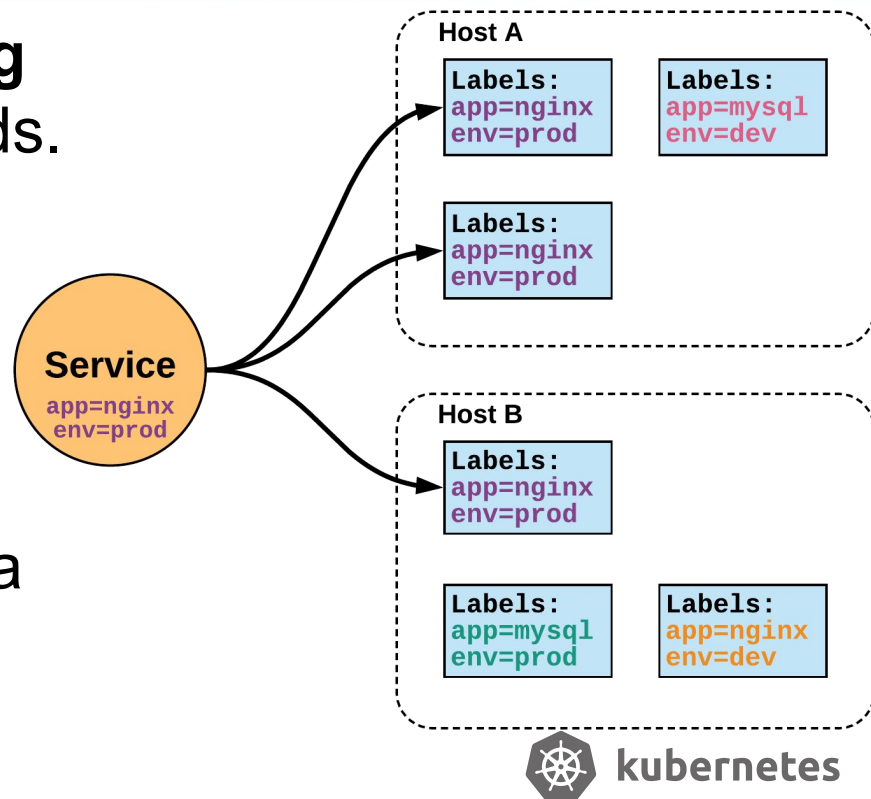
**Set-based** selectors are supported on a limited subset of objects. However, they provide a method of filtering on a set of values, and supports multiple operators including: **in**, **notin**, and **exist**.

```
selector:  
  matchExpressions:  
    - key: gpu  
      operator: in  
      values: ["nvidia"]
```



## 4. Services

- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource**
  - static cluster IP
  - static namespaced DNS name
- An abstraction which defines a logical set of Pods
- Uses pod readiness probes



# Probes



- **Readiness Probe:**

- Indicates whether the Container is ready to service requests. Cont. may take time to load data etc.
- Used by Services

- **Liveness Probe:**

- Indicates whether the Container is running.
- If it fails, kubernetes kills the Container
- Useful to detect app deadlocks

# Readiness Probe



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    resources: {}
    readinessProbe:
      httpGet:
        path: /healthz
        port: 80
      initialDelaySeconds: 100
      periodSeconds: 10
      timeoutSeconds: 4
      failureThreshold: 2
```



# Liveness Probe



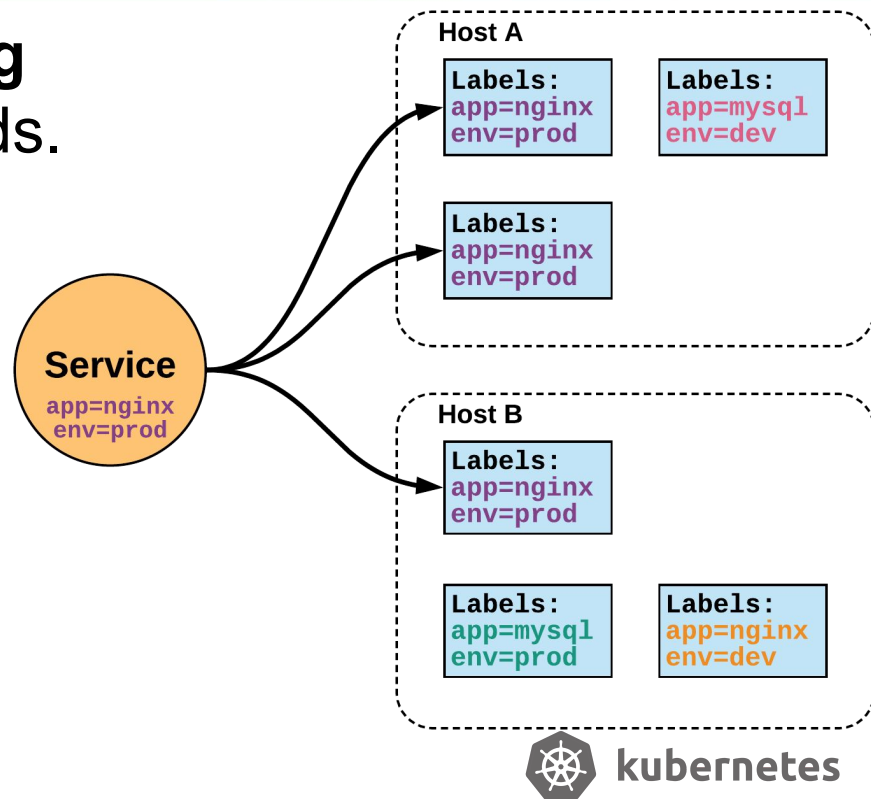
```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
  livenessProbe:
    httpGet:
      path: /healthz
      port: 80
    initialDelaySeconds: 100
    periodSeconds: 10
    timeoutSeconds: 4
    failureThreshold: 2
```



# 3. Services

- **Unified method of accessing** the exposed workloads of Pods.
- **Durable resource**
  - static cluster IP
  - static namespaced DNS name

**NOT Ephemeral !!!**





## 3. Services



There are 4 major service types:

- **ClusterIP** (default)
- **NodePort**
- **LoadBalancer**
- **ExternalName**

# ClusterIP Service



**ClusterIP** services exposes a service on a strictly cluster internal virtual IP.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  selector:
    app: nginx
    env: prod
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

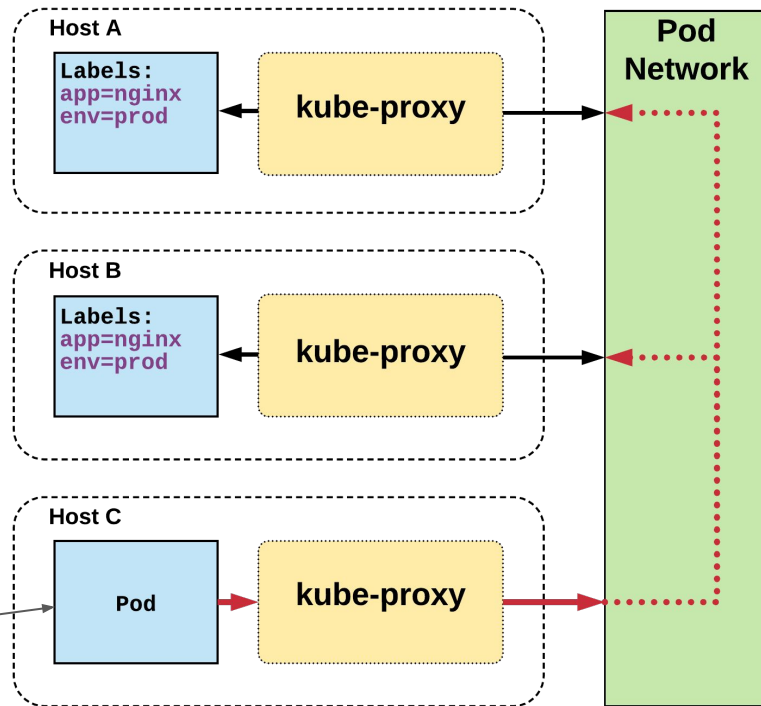


# ClusterIP Service



Name: example-prod  
Selector: app=nginx,env=prod  
Type: ClusterIP  
IP: 10.96.28.176  
Port: <unset> 80/TCP  
TargetPort: 80/TCP  
Endpoints: 10.255.16.3:80,  
10.255.16.4:80

```
/ # nslookup example-prod.default.svc.cluster.local  
Name:      example-prod.default.svc.cluster.local  
Address 1: 10.96.28.176 example-prod.default.svc.cluster.local
```



kubernetes

# Node Port Service

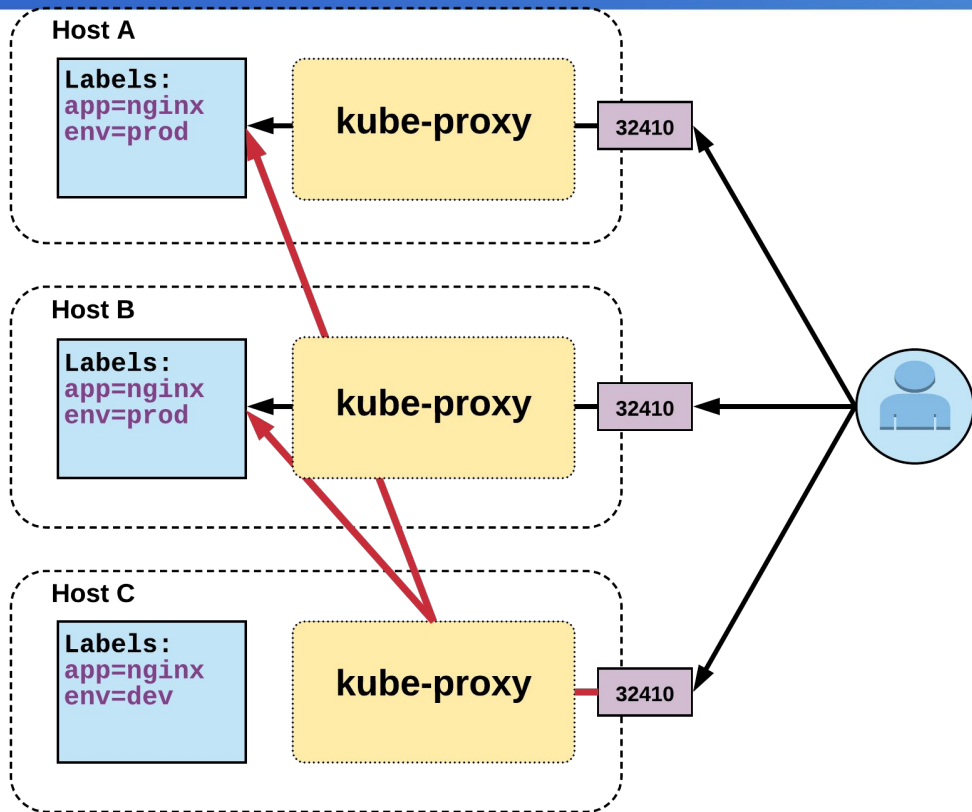


- **NodePort** services extend the **ClusterIP** service.
- Exposes a port on every node's IP.
- Port can either be statically defined, or dynamically taken from a range between 30000-32767.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: NodePort
  selector:
    app: nginx
    env: prod
  ports:
    - nodePort: 32410
      protocol: TCP
      port: 80
      targetPort: 80
```



# Node Port Service



Name: example-prod  
Selector: app=nginx,env=prod  
Type: NodePort  
IP: 10.96.28.176  
Port: <unset> 80/TCP  
TargetPort: 80/TCP  
NodePort: <unset> 32410/TCP  
Endpoints: 10.255.16.3:80,  
10.255.16.4:80



kubernetes

# Load Balancer Service

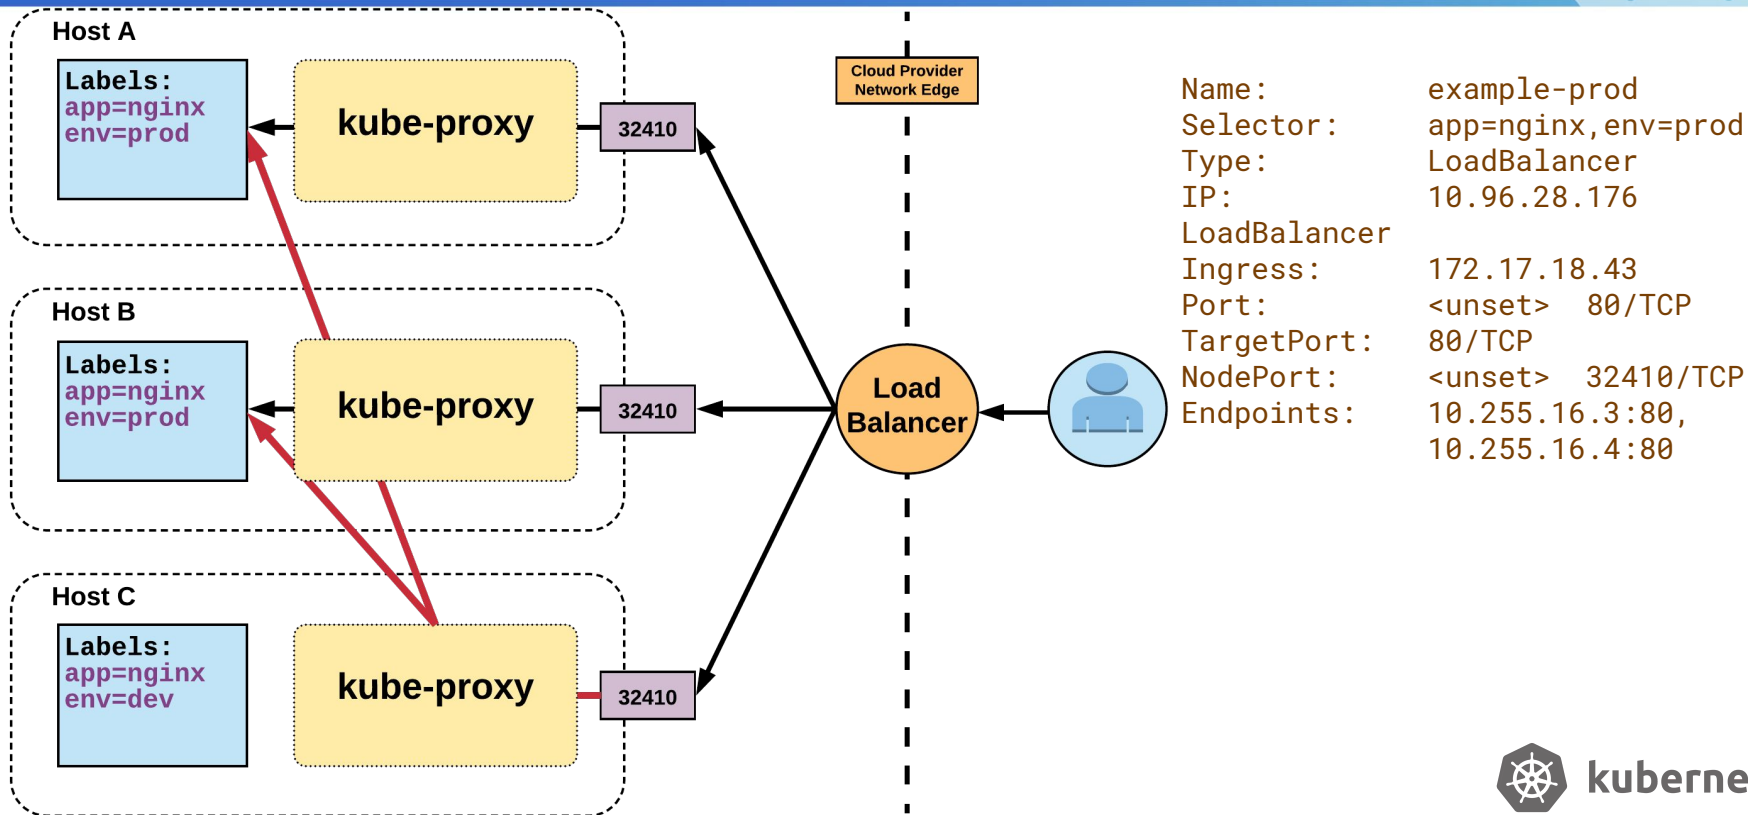


- **LoadBalancer** services extend **NodePort**.
- Works in conjunction with an external system to map a cluster external IP to the exposed service.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: LoadBalancer
  selector:
    app: nginx
    env: prod
  ports:
    protocol: TCP
    port: 80
    targetPort: 80
```



# Load Balancer Service



# ExternalName Service



- **ExternalName** is used to reference endpoints **OUTSIDE** the cluster.
- Creates an internal **CNAME** DNS entry that aliases another.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: ExternalName
spec:
  externalName:
    example.com
```





## 4. ReplicaSets



- Primary method of managing pod replicas and their lifecycle.
- Includes their scheduling, scaling, and deletion.
- Their job is simple: **Always ensure the desired number of pods are running.**



## 4. ReplicaSets



- **replicas**: The desired number of instances of the Pod.
- **selector**: The label selector for the **ReplicaSet** will manage **ALL** Pod instances that it targets; whether it's desired or not.

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-example
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  template:
    <pod template>
```




## 4. ReplicaSets



```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: rs-example
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  template:
    metadata:
      labels:
        app: nginx
        env: prod
    spec:
      containers:
        - name: nginx
          image: nginx:stable-alpine
          ports:
            - containerPort: 80
```

# Architecture



# Control Plane Components



- Kube API Server
- ETCD
- Kube Controller Manager
- Kube-Scheduler

# Kube API Server



- Provides a forward facing REST interface into the kubernetes control plane and datastore.
- All clients and other applications interact with kubernetes **strictly** through the API Server.



# ETCD



- ETCD acts as the cluster datastore.
- Purpose in relation to Kubernetes is to provide a strong, consistent and highly available key-value store for persisting cluster state.
- Stores objects and config information.



# Kube Controller Manager



- Monitors the cluster state via the API server and **steers the cluster towards the desired state**. (Reconciliation loop - Imperative vs Declarative)
- **Node Controller**: Responsible for noticing and responding when nodes go down.
- **Replication Controller**: Responsible for maintaining the correct number of pods for every replication controller object in the system.
- **Endpoints Controller**: Populates the Endpoints object (that is, joins Services & Pods).



# Kube Scheduler



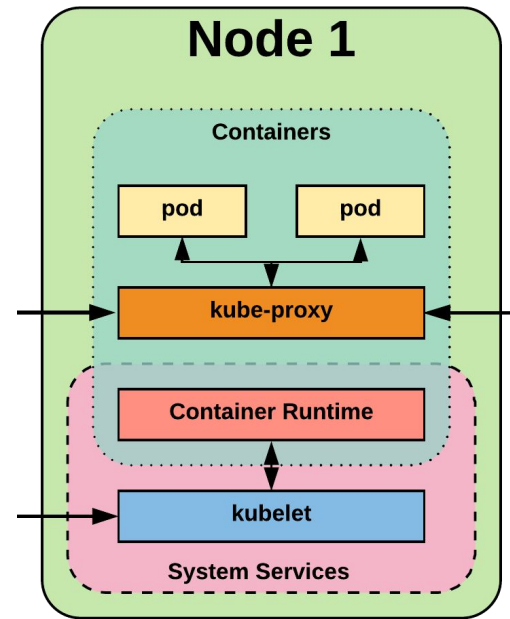
- Component on the master that **watches newly created pods** that have no node assigned, and **selects a node for them** to run on.
- Factors taken into account for scheduling decisions include individual and collective resource requirements, hardware/software/policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference and deadlines.
- Can be **extended**.



# Node Components



- Kubelet
- KubeProxy
- Container Runtime Engine



kubernetes

# Kubelet



- An **agent** that runs on each node. It makes sure that containers are running in a pod.
- **Takes a set of PodSpecs** and ensures that the containers described in those PodSpecs are running and healthy.
- **On workers:** poll kube-apiserver looking for what they should run
- **On masters:** run the master services as static manifests found locally on the host (/etc/kubernetes/manifests/)

# KubeProxy




- Manages the network rules on each node.
- Performs connection forwarding or load balancing for Kubernetes cluster services.

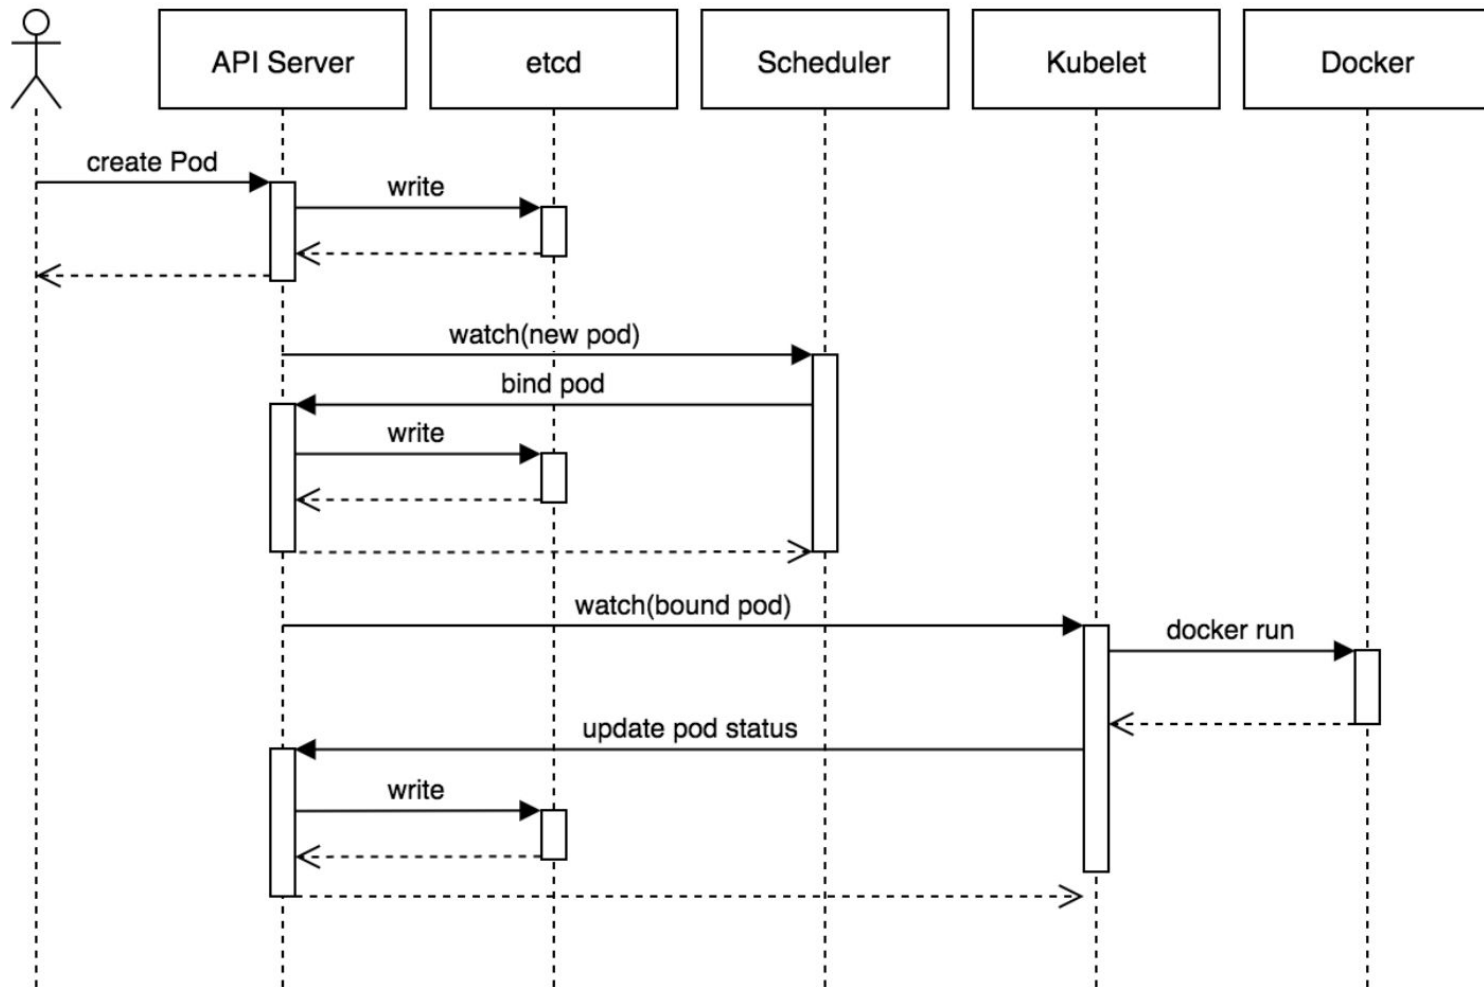
# Container Runtime Engine



- A container runtime is a CRI (Container Runtime Interface) compatible application that executes and manages containers.
  - Containerd (docker)
  - CRI-o
  - Rkt
  - Kata (formerly clear and hyper)
  - Virtlet (VM CRI compatible runtime)



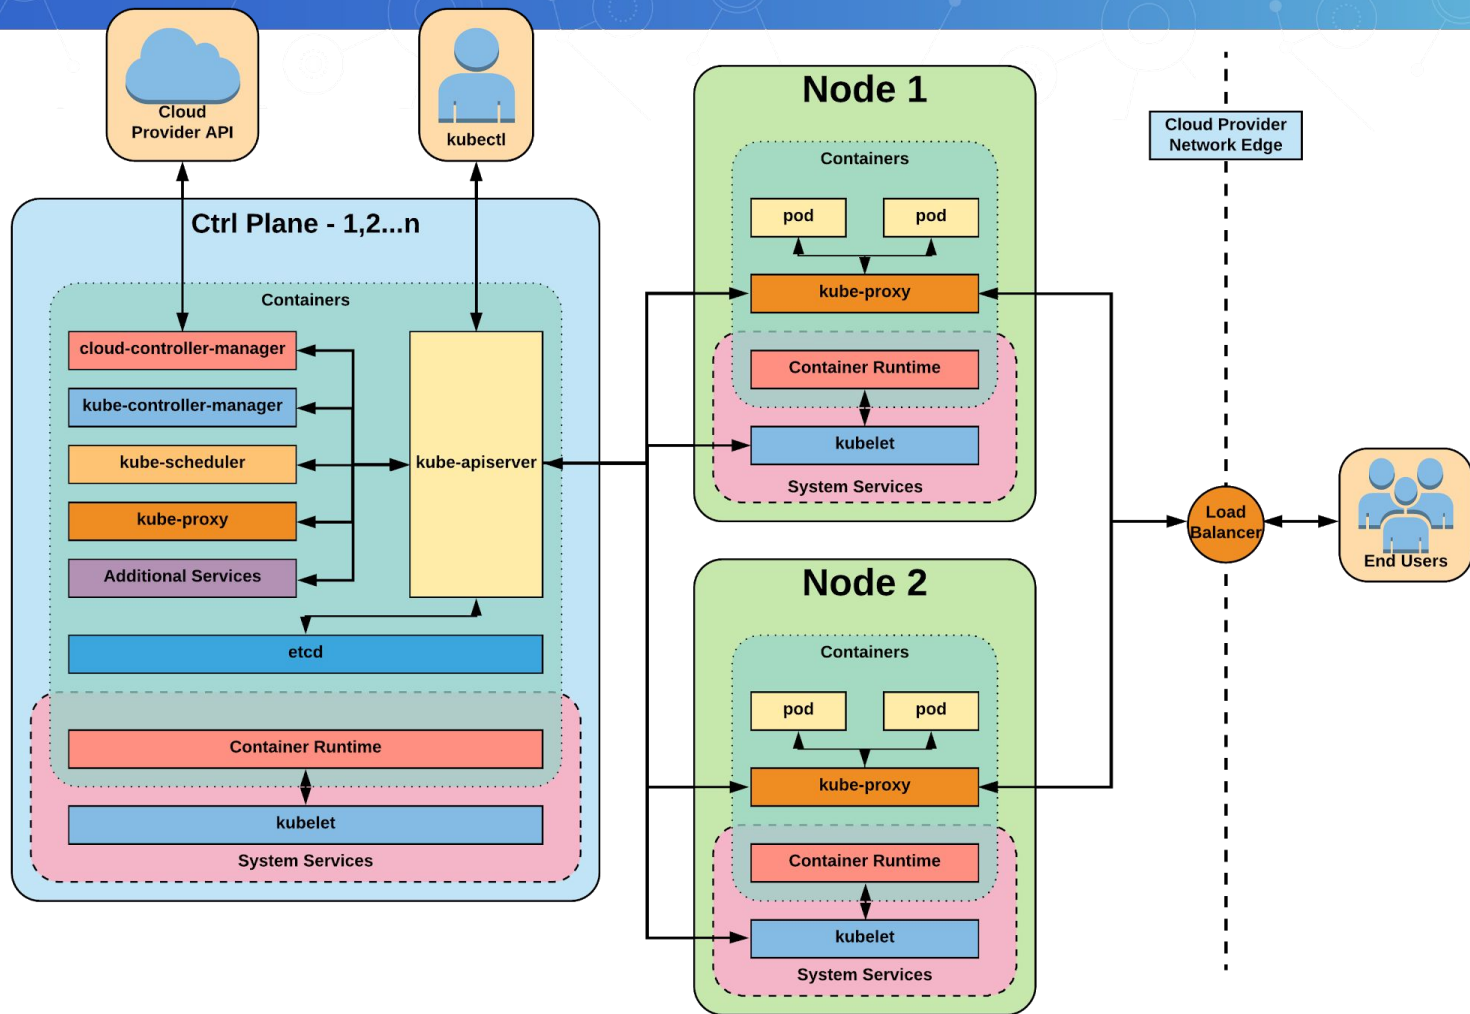




# Summery








# Q&A

- by Joe Beda (Gluecon 2017)

This presentation is licensed under a Creative Commons Attribution 4.0 International License.

See <https://creativecommons.org/licenses/by/4.0/> for more details.



# Patterns for composite containers



- Sidecar
- Ambassador
- Adaptor

<https://kubernetes.io/blog/2015/06/the-distributed-system-toolkit-patterns/>

<https://storage.googleapis.com/pub-tools-public-publication-data/pdf/45406.pdf>



kubernetes